

SQL: The Query Language

R & G - Chapter 5



Relational Tables



- *Schema* is fixed:
 - attribute names, *atomic* types
 - students(name text, gpa float, dept text)
- *Instance* can change
 - a *multiset* of “rows” (“tuples”)
 - {(‘Bob Snob’, 3.3, ‘CS’),
(‘Bob Snob’, 3.3, ‘CS’),
(‘Mary Contrary’, 3.8, ‘CS’)}

Basic Single-Table Queries



```
SELECT [DISTINCT] <column expression list>
FROM <single table>
[WHERE <predicate>]
[GROUP BY <column list>]
[HAVING <predicate>] ]
[ORDER BY <column list>];
```

Basic Single-Table Queries



```
SELECT [DISTINCT] <column expression list>
FROM <single table>
[WHERE <predicate>]
[GROUP BY <column list>]
[HAVING <predicate>] ]
[ORDER BY <column list>];
```

Simplest version is straightforward

- Produce all tuples in the table that satisfy the predicate
- Output the expressions in the SELECT list
- Expression can be a column reference, or an arithmetic expression over column refs

Basic Single-Table Queries



```
SELECT S.name, S.gpa
FROM students AS S
WHERE S.dept = 'CS'
[GROUP BY <column list>]
[HAVING <predicate>] ]
[ORDER BY <column list>];
```

Simplest version is straightforward

- Produce all tuples in the table that satisfy the predicate
- Output the expressions in the SELECT list
- Expression can be a column reference, or an arithmetic expression over column refs

SELECT DISTINCT



```
SELECT DISTINCT S.name, S.gpa
FROM students S
WHERE S.dept = 'CS'
[GROUP BY <column list>]
[HAVING <predicate>] ]
[ORDER BY <column list>];
```

DISTINCT flag specifies removal of duplicates before output

Removed the “AS” from FROM clause --- it’s optional

ORDER BY



```
SELECT DISTINCT S.name, S.gpa, S.age*2 AS a2
FROM Students S
WHERE S.dept = 'CS'
[GROUP BY <column list>
[HAVING <predicate>] ]
ORDER BY S.gpa, S.name, a2;
```

ORDER BY clause specifies output to be sorted

- *Lexicographic* ordering (left to right)

Obviously must refer to columns in the output

- Note the AS clause for naming output columns!

ORDER BY



```
SELECT DISTINCT S.name, S.gpa
FROM Students S
WHERE S.dept = 'CS'
[GROUP BY <column list>
[HAVING <predicate>] ]
ORDER BY S.gpa DESC, S.name ASC;
```

Ascending order by default, but can be overridden

- DESC flag for descending, ASC for ascending
- Can mix and match, lexicographically

AGGREGATES



```
SELECT [DISTINCT] AVG(S.gpa)
FROM Students S
WHERE S.dept = 'CS'
[GROUP BY <column list>
[HAVING <predicate>] ]
[ORDER BY <column list>] ;
```

Before producing output, compute a summary
(a.k.a. an *aggregate*) of some arithmetic expression

Produces 1 row of output

- with one column in this case

Other aggregates: SUM, COUNT, MAX, MIN

Note: can use DISTINCT *inside* the agg function

- SELECT COUNT(DISTINCT S.name) FROM Students S
- vs. SELECT DISTINCT COUNT (S.name) FROM Students S;

GROUP BY



```
SELECT [DISTINCT] AVG(S.gpa), S.dept
FROM Students S
[WHERE <predicate>]
GROUP BY S.dept
[HAVING <predicate>]
[ORDER BY <column list>] ;
```

Partition table into groups with same GROUP BY column values

- Can group by a list of columns

Produce an aggregate result per group

- Cardinality of output = # of distinct group values

Note: can put grouping columns in SELECT list

- For aggregate queries, SELECT list can contain aggs and GROUP BY columns only!
- What would it mean if we said SELECT S.name, AVG(S.gpa) above??

HAVING



```
SELECT [DISTINCT] AVG(S.gpa), S.dept
FROM Students S
[WHERE <predicate>]
GROUP BY S.dept
HAVING COUNT(*) > 5
[ORDER BY <column list>] ;
```

The HAVING predicate is applied *after* grouping and aggregation

- Hence can contain anything that could go in the SELECT list
- I.e. aggs or GROUP BY columns

HAVING can only be used in aggregate queries

(It's an optional clause for GROUP BY)

Putting it all together



```
SELECT S.dept, AVG(S.gpa), COUNT(*)
FROM Students S
WHERE S.gender = 'F'
GROUP BY S.dept
HAVING COUNT(*) > 5
ORDER BY S.dept;
```

Relational Query Languages



Two sublanguages:

- DDL – Data Definition Language
 - Define and modify schema
- DML – Data Manipulation Language
 - Write declarative queries/updates
 - We just covered basic queries is the SQL DML

DBMS is responsible for efficient evaluation

- Semantics are precise (more on that later)
- Declarative language => room for optimization
- Optimizer can re-order operations
 - Won't affect query answer

Example Database



Sailors

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

Boats

bid	bname	color
101	Nina	red
102	Pinta	blue
103	Santa Maria	red

Reserves

sid	bid	day
1	102	9/12/2015
2	102	9/13/2015

SQL DDL



```
CREATE TABLE Sailors (
  sid INTEGER,
  sname CHAR(20),
  rating INTEGER,
  age REAL,
  PRIMARY KEY (sid));
```

```
CREATE TABLE Boats (
  bid INTEGER,
  bname CHAR(20),
  color CHAR(10),
  PRIMARY KEY (bid));
```

```
CREATE TABLE Reserves (
  sid INTEGER,
  bid INTEGER,
  day DATE,
  PRIMARY KEY (sid, bid, day),
  FOREIGN KEY (sid) REFERENCES Sailors,
  FOREIGN KEY (bid) REFERENCES Boats);
```

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

bid	bname	color
101	Nina	red
102	Pinta	blue
103	Santa Maria	red

sid	bid	day
1	102	9/12
2	102	9/13

Querying Multiple Relations



```
SELECT S.sname
FROM   Sailors AS S, Reserves AS R
WHERE  S.sid=R.sid AND R.bid=102
```

Sailors

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

Reserves

sid	bid	day
1	102	9/12
2	102	9/13

Conceptual SQL Evaluation




```
SELECT  [DISTINCT] target-list
FROM    relation-list
WHERE   qualification
GROUP BY grouping-list
HAVING  group-qualification
```

Project away columns
(just keep those used in
SELECT, GROUP BY,
HAVING)

Apply selections
(eliminate rows)

Relation
cross-product



Query Semantics

```
SELECT [DISTINCT] target-list
FROM   relation-list
WHERE  qualification
```

1. FROM : compute cross product of tables.
2. WHERE : Check conditions, discard tuples that fail.
3. SELECT : Delete unwanted fields.
4. DISTINCT (optional) : eliminate duplicate rows.

Note: likely a terribly inefficient strategy!

- Query optimizer will find more efficient plans.

Find sailors who've reserved at least one boat

```
SELECT S.sid
FROM   Sailors AS S, Reserves AS R
WHERE  S.sid=R.sid
```

Would DISTINCT make a difference here?

About Range Variables

Needed when ambiguity could arise.

- e.g., same table used multiple times in FROM (“self-join”)

```
SELECT x.sname, x.age, y.sname, y.age
FROM   Sailors AS x, Sailors AS y
WHERE  x.age > y.age
```

Sailors

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

Arithmetic Expressions

```
SELECT S.age, S.age-5 AS age1, 2*S.age AS age2
FROM   Sailors AS S
WHERE  S.sname = 'dustin'
```

```
SELECT S1.sname AS name1, S2.sname AS name2
FROM   Sailors AS S1, Sailors AS S2
WHERE  2*S1.rating = S2.rating - 1
```

String Comparisons

```
SELECT S.sname
FROM   Sailors S
WHERE  S.sname LIKE 'B_%B'
```

‘_’ stands for any one character and ‘%’ stands for 0 or more arbitrary characters.

Most DBMSs now support standard regex as well

Find sid's of sailors who've reserved a red **or** a green boat

```
SELECT R.sid
FROM   Boats B, Reserves R
WHERE  R.bid=B.bid AND
       (B.color='red' OR
        B.color='green')
```

... OR:

```
SELECT R.sid
FROM   Boats B, Reserves R
WHERE  R.bid=B.bid AND
       B.color='red'
UNION
SELECT R.sid
FROM   Boats B, Reserves R
WHERE  R.bid=B.bid AND B.color='green'
```

Find sid's of sailors who've reserved a red **and** a green boat

```
SELECT R.sid
FROM   Boats B, Reserves R
WHERE  R.bid=B.bid AND
       (B.color='red' AND B.color='green')
```

Find sid's of sailors who've reserved a red **and** a green boat



```
SELECT S.sid
FROM   Sailors S, Boats B, Reserves R
WHERE  S.sid=R.sid
        AND R.bid=B.bid
        AND B.color='red'

INTERSECT

SELECT S.sid
FROM   Sailors S, Boats B, Reserves R
WHERE  S.sid=R.sid
        AND R.bid=B.bid
        AND B.color='green'
```

Two sets must match in columns/types
Whole tuple must match

Find sid's of sailors who've reserved a red **and** a green boat



Could alternatively use a self-join:

```
SELECT R1.sid
FROM   Boats B1, Reserves R1,
        Boats B2, Reserves R2
WHERE  R1.sid=R2.sid
        AND R1.bid=B1.bid
        AND R2.bid=B2.bid
        AND (B1.color='red' AND B2.color='green')
```

Find sid's of sailors who have not reserved a boat



```
SELECT S.sid
FROM   Sailors S

EXCEPT

SELECT S.sid
FROM   Sailors S, Reserves R
WHERE  S.sid=R.sid
```

Nested Queries: IN



Names of sailors who've reserved boat #102:

```
SELECT S.sname
FROM   Sailors S
WHERE  S.sid IN
        (SELECT R.sid
         FROM   Reserves R
         WHERE  R.bid=102)
```

Nested Queries: NOT IN



Names of sailors who've not reserved boat #103:

```
SELECT S.sname
FROM   Sailors S
WHERE  S.sid NOT IN
        (SELECT R.sid
         FROM   Reserves R
         WHERE  R.bid=103)
```

Nested Queries with Correlation



Names of sailors who've reserved boat #102:

```
SELECT S.sname
FROM   Sailors S
WHERE  EXISTS
        (SELECT *
         FROM   Reserves R
         WHERE  R.bid=102 AND S.sid=R.sid)
```

- Subquery must be recomputed for each Sailors tuple.
 - Think of subquery as a function call that runs a query

More on Set-Comparison Operators



- we've seen: IN, EXISTS
- can also have: NOT IN, NOT EXISTS
- other forms: op ANY, op ALL

Find sailors whose rating is greater than that of *some* sailor called Fred:

```
SELECT *
FROM Sailors S
WHERE S.rating > ANY
  (SELECT S2.rating
   FROM Sailors S2
   WHERE S2.sname='Fred')
```

A Tough One



Find sailors who've reserved all boats.

```
SELECT S.sname   Sailors S such that ...
FROM Sailors S
WHERE NOT EXISTS (SELECT B.bid   there is no boat B
                  FROM Boats B   without ...
                  WHERE NOT EXISTS (SELECT R.bid
                                    FROM Reserves R
                                    WHERE R.bid=B.bid
                                    AND R.sid=S.sid ))
a Reserves tuple showing S reserved B
```

ARGMAX?



The sailor with the highest rating
– what about ties for highest?!

```
SELECT *
FROM Sailors S
WHERE S.rating >= ALL
  (SELECT S2.rating
   FROM Sailors S2)
```

```
SELECT *
FROM Sailors S
WHERE S.rating =
  (SELECT MAX(S2.rating)
   FROM Sailors S2)
```

```
SELECT *
FROM Sailors S
ORDER BY rating DESC
LIMIT 1;
```

Null Values



Field values are sometimes unknown or inapplicable
– SQL provides a special value **null** for such situations.

The presence of null complicates many issues. E.g.:

- Special syntax "IS NULL" and "IS NOT NULL"
- Assume rating = NULL. Consider predicate "rating>8".
 - True? False? (answer is always false)
 - What about AND, OR and NOT connectives?
 - SUM?
- We need a 3-valued logic (true, false and unknown).
- Meaning of constructs must be defined carefully. (e.g., WHERE clause eliminates rows that don't evaluate to true.)
- New operators (in particular, **outer joins**) possible/needed.