

## SQL: The Query Language (part 2 of 2)

R & G - Chapter 5



## Basic Single-Table Queries



```
SELECT [DISTINCT] <column expression list>
FROM <single table>
[WHERE <predicate>]
[GROUP BY <column list>]
[HAVING <predicate>] ]
[ORDER BY <column list>];
```

## Null Values



Field values are sometimes unknown or inapplicable  
– SQL provides a special value **null** for such situations.

The presence of null complicates many issues. E.g.:

- Special syntax “IS NULL” and “IS NOT NULL”
- Assume rating = NULL. Consider predicate “rating>8”.
  - True? False? (answer is always *false*)
  - What about AND, OR and NOT connectives?
  - SUM?
- We need a 3-valued logic (true, false and unknown).
- Meaning of constructs must be defined carefully. (e.g., WHERE clause eliminates rows that don't evaluate to true.)
- New operators (in particular, **outer joins**) possible/needed.

## Joins



```
SELECT (column_list)
FROM table_name
[INNER | {LEFT | RIGHT | FULL} {OUTER}] JOIN table_name
ON qualification_list
WHERE ...
```

INNER is default

## Inner/Natural Joins



```
SELECT s.sid, s.sname, r.bid
FROM Sailors s, Reserves r
WHERE s.sid = r.sid
```

```
SELECT s.sid, s.sname, r.bid
FROM Sailors s INNER JOIN Reserves r
ON s.sid = r.sid
```

all 3 are  
equivalent!

```
SELECT s.sid, s.sname, r.bid
FROM Sailors s NATURAL JOIN Reserves r
```

“NATURAL” means equi-join for each pair of attributes with the same name

```
SELECT s.sid, s.sname, r.bid
FROM Sailors2 s INNER JOIN Reserves2 r
ON s.sid = r.sid;
```



sid	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
95	Bob	3	63.5

sid	bid	day
22	101	10/10/96
95	103	11/12/96

s.sid	s.name	r.bid
22	Dustin	101
95	Bob	103

## Left Outer Join



Returns all matched rows, plus all unmatched rows from the table on the left of the join clause  
(use nulls in fields of non-matching tuples)

```
SELECT s.sid, s.sname, r.bid
FROM Sailors2 s LEFT OUTER JOIN Reserves2 r
ON s.sid = r.sid;
```

Returns all sailors & bid for boat in any of their reservations

Note: no match for s.sid? => r.bid is NULL

```
SELECT s.sid, s.sname, r.bid
FROM Sailors2 s LEFT OUTER JOIN Reserves2 r
ON s.sid = r.sid
```



sid	sname	rating	age	sid	bid	day
22	Dustin	7	45.0	22	101	10/10/96
31	Lubber	8	55.5	95	103	11/12/96
95	Bob	3	63.5			

s.sid	s.sname	r.bid
22	Dustin	101
95	Bob	103
31	Lubber	

## Right Outer Join



Right Outer Join returns all matched rows, plus all unmatched rows from the table on the right of the join clause

```
SELECT r.sid, b.bid, b.bname
FROM Reserves2 r RIGHT OUTER JOIN Boats2 b
ON r.bid = b.bid;
```

Returns all boats & information on which are reserved

No match for b.bid? => r.sid is NULL

```
SELECT r.sid, b.bid, b.bname
FROM Reserves2 r RIGHT OUTER JOIN Boats2 b
ON r.bid = b.bid;
```



sid	bid	day	bid	bname	color
22	101	10/10/96	101	Interlake	blue
95	103	11/12/96	102	Interlake	red
			103	Clipper	green
			104	Marine	red

r.sid	b.bid	b.name
22	101	Interlake
	102	Interlake
95	103	Clipper
	104	Marine

## Full Outer Join



Full Outer Join returns all (matched or unmatched) rows from the tables on both sides of the join clause

```
SELECT r.sid, b.bid, b.bname
FROM Reserves2 r FULL OUTER JOIN Boats2 b
ON r.bid = b.bid
```

- Returns all boats & all reservations
- No match for r.bid?
  - b.bid is NULL AND b.bname is NULL
- No match for b.bid?
  - r.sid is NULL

```
SELECT r.sid, b.bid, b.bname
FROM Reserves2 r FULL OUTER JOIN Boats2 b
ON r.bid = b.bid
```



sid	bid	day	bid	bname	color
22	101	10/10/96	101	Interlake	blue
95	103	11/12/96	102	Interlake	red
			103	Clipper	green
			104	Marine	red

r.sid	b.bid	b.name
22	101	Interlake
	102	Interlake
95	103	Clipper
	104	Marine

**Note: in this case it is the same as the ROJ!**  
*bid* is a foreign key in reserves, so all reservations must have a corresponding tuple in boats

## Views: Named Queries



```
CREATE VIEW view_name
AS select_statement
```

Makes development simpler  
Often used for security  
(not “materialized”)

```
CREATE VIEW Reds
AS SELECT B.bid, COUNT (*) AS scount
FROM Boats2 B, Reserves2 R
WHERE R.bid=B.bid AND B.color='red'
GROUP BY B.bid
```

## Views Instead of Relations in Queries



```
CREATE VIEW Redcount
AS SELECT B.bid, COUNT (*) AS scount
FROM Boats2 B, Reserves2 R
WHERE R.bid=B.bid AND B.color='red'
GROUP BY B.bid
```

bid	scount
102	1

Reds

```
SELECT bname, scount
FROM Redcount R, Boats2 B
WHERE R.bid=B.bid
AND scount < 10
```

## Subqueries in FROM



```
SELECT bname, scount
FROM Boats2 B,
  (SELECT B.bid, COUNT (*)
   FROM Boats2 B, Reserves2 R
   WHERE R.bid=B.bid AND B.color='red'
   GROUP BY B.bid) AS Reds(bid, scount)
WHERE Reds.bid=B.bid
AND scount < 10
```

## WITH (common table expression)



```
WITH Reds(bid, scount) AS
  (SELECT B.bid, COUNT (*)
   FROM Boats2 B, Reserves2 R
   WHERE R.bid=B.bid AND B.color='red'
   GROUP BY B.bid)
SELECT bname, scount
FROM Boats2 B, Reds
WHERE Reds.bid=B.bid
AND scount < 10
```

## Discretionary Access Control



```
GRANT privileges ON object TO users
[WITH GRANT OPTION]
```

- Object can be a **Table** or a **View**
- Privileges can be:
  - Select
  - Insert
  - Delete
  - References (cols) – allow to create a foreign key that references the specified column(s)
  - All
- Can later be REVOKED
- Users can be single users or groups
- See Chapter 17 for more details.

## Two more important topics



- Constraints
- SQL embedded in other languages

## Integrity Constraints



- IC conditions that every legal instance of a relation must satisfy.
  - Inserts/deletes/updates that violate ICs are disallowed.
  - Can ensure application semantics (e.g., sid is a key),
  - ...or prevent inconsistencies (e.g., sname has to be a string, age must be < 200)
- Types of IC's: Domain constraints, primary key constraints, foreign key constraints, general constraints.
  - Domain constraints: Field values must be of right type. Always enforced.
  - Primary key and foreign key constraints: coming right up.

## Where do ICs Come From?




- Semantics of the real world!
- Note:
  - We can check IC violation in a DB instance
  - We can NEVER infer that an IC is true by looking at an instance.
    - An IC is a statement about all possible instances!
  - From example, we know name is not a key, but the assertion that sid is a key is given to us.
- Key and foreign key ICs are the most common
- More general ICs supported too.

## Keys



- Keys are a way to associate tuples in different relations
- Keys are one form of IC



## Primary Keys



- A set of fields is a **superkey** if:
  - No two distinct tuples can have same values in all key fields
- A set of fields is a **key** for a relation if it is *minimal*:
  - It is a superkey
  - No subset of the fields is a superkey
- what if >1 key for a relation?
  - One of the keys is chosen (by DBA) to be the **primary key**. Other keys are called **candidate keys**.
- E.g.
  - sid is a key for Students.
  - What about name?
  - The set {sid, gpa} is a superkey.

## Primary and Candidate Keys



- Possibly many **candidate keys** (specified using **UNIQUE**), one of which is chosen as the **primary key**.
- Keys must be used carefully!

```
CREATE TABLE Enrolled1 (sid CHAR(20), cid CHAR(20), grade CHAR(2), PRIMARY KEY (sid,cid))
CREATE TABLE Enrolled2 (sid CHAR(20), cid CHAR(20), grade CHAR(2), UNIQUE (cid, grade))
```

vs.

"For a given student and course, there is a single grade."

## Primary and Candidate Keys



```
CREATE TABLE Enrolled1 (sid CHAR(20), cid CHAR(20), grade CHAR(2), PRIMARY KEY (sid,cid))
CREATE TABLE Enrolled2 (sid CHAR(20), cid CHAR(20), grade CHAR(2), PRIMARY KEY (sid), UNIQUE (cid, grade))
```

vs.

```
INSERT INTO enrolled1 VALUES ('1234', 'cs186', 'A+');
INSERT INTO enrolled1 VALUES ('1234', 'cs186', 'F');
INSERT INTO enrolled1 VALUES ('1234', 'cs61C', 'A+');
```

```
INSERT INTO enrolled2 VALUES ('1234', 'cs186', 'A+');
INSERT INTO enrolled2 VALUES ('1234', 'cs186', 'F');
INSERT INTO enrolled2 VALUES ('1234', 'cs61C', 'A+');
INSERT INTO enrolled2 VALUES ('4567', 'cs186', 'A+');
```

"For a given student and course, there is a single grade."

## Primary and Candidate Keys



```
CREATE TABLE Enrolled1
(sid CHAR(20),
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid,cid));
```

VS.

```
CREATE TABLE Enrolled2
(sid CHAR(20),
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid),
UNIQUE (cid, grade));
```

```
INSERT INTO enrolled1 VALUES ('1234', 'cs186', 'A+');
INSERT INTO enrolled1 VALUES ('1234', 'cs186', 'F');
INSERT INTO enrolled1 VALUES ('1234', 'cs61C', 'A+');
```

```
INSERT INTO enrolled2 VALUES ('1234', 'cs186', 'A+');
INSERT INTO enrolled2 VALUES ('1234', 'cs186', 'F');
INSERT INTO enrolled2 VALUES ('1234', 'cs61C', 'A+');
INSERT INTO enrolled2 VALUES ('4567', 'cs186', 'A+');
```

*"Students can take only one course, and no two students in a course receive the same grade."*

## Foreign Keys, Referential Integrity



- **Foreign key:** a "logical pointer"
  - Set of fields in a tuple in one relation that 'refer' to a tuple in another relation.
  - Reference to *primary key* of the other relation.
- All foreign key constraints enforced?
  - **referential integrity!**
  - i.e., no dangling references.

## Foreign Keys in SQL



- E.g. Only students listed in the Students relation should be allowed to enroll for courses.
  - *sid* is a foreign key referring to **Students**:

```
CREATE TABLE Enrolled
(sid CHAR(20),cid CHAR(20),grade CHAR(2),
PRIMARY KEY (sid,cid),
FOREIGN KEY (sid) REFERENCES Students);
```

Enrolled

sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B
<del>11111</del>	<del>English102</del>	<del>A</del>

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

## Enforcing Referential Integrity



- *sid* in Enrolled: foreign key referencing Students.
- Scenarios:
  - Insert Enrolled tuple with non-existent student id?
  - Delete a Students tuple?
    - Also delete Enrolled tuples that refer to it? (CASCADE)
    - Disallow if referred to? (NO ACTION)
    - Set *sid* in referring Enrolled tups to a *default* value? (SET DEFAULT)
    - Set *sid* in referring Enrolled tuples to *null*, denoting 'unknown' or 'inapplicable'. (SET NULL)
- Similar issues arise if primary key of Students tuple is updated.

## General Constraints

```
CREATE TABLE Sailors
(sid INTEGER,
sname CHAR(10),
rating INTEGER,
age REAL,
PRIMARY KEY (sid),
CHECK (rating >= 1
AND rating <= 10 ))
```



- Useful when more general ICs than keys are involved.
- Can use queries to express constraint.
- Checked on insert or update.
- Constraints can be named.

```
CREATE TABLE Reserves
(sname CHAR(10),
bid INTEGER,
day DATE,
PRIMARY KEY (bid,day),
CONSTRAINT noInterlakeRes
CHECK ('Interlake' <>
(SELECT B.bname
FROM Boats B
WHERE B.bid=bid)))
```

## Constraints Over Multiple Relations



```
CREATE TABLE Sailors
(sid INTEGER,
sname CHAR(10),
rating INTEGER,
age REAL,
PRIMARY KEY (sid),
CHECK
( (SELECT COUNT (S.sid) FROM Sailors S)
+ (SELECT COUNT (B.bid) FROM
Boats B) < 100 ))
```

Number of boats  
plus number of  
sailors is < 100

## Constraints Over Multiple Relations



- Awkward and wrong!
  - Only checks sailors!
- ASSERTION is the right solution; not associated with either table.

- Unfortunately, not supported in many DBMS.
- Triggers are another solution.

```
CREATE TABLE Sailors
(sid INTEGER,
sname CHAR(10),
rating INTEGER,
age REAL,
PRIMARY KEY (sid),
CHECK
( (SELECT COUNT (S.sid) FROM Sailors S)
+ (SELECT COUNT (B.bid) FROM
Boats B) < 100 )
```

Number of boats  
plus number of  
sailors is < 100

```
CREATE ASSERTION smallClub
CHECK
( (SELECT COUNT (S.sid) FROM Sailors S)
+ (SELECT COUNT (B.bid)
FROM Boats B) < 100 )
```

## Two more important topics



- **Constraints**
- SQL embedded in other languages

## Writing Applications with SQL



- SQL is not a general purpose programming language.
  - + Tailored for data retrieval and manipulation
  - + Relatively easy to optimize and parallelize
  - Can't write entire apps in SQL alone
- Options:
  - Make the query language "Turing complete"
    - Avoids the "impedance mismatch"
    - makes "simple" relational language complex
  - Allow SQL to be embedded in regular programming languages.
  - Q: What needs to be solved to make the latter approach work?

## Cursors



- Can declare a cursor on a relation or query
- Can *open* a cursor
- Can repeatedly *fetch* a tuple (moving the cursor)
- Special return value when all tuples have been retrieved.
- **ORDER BY** allows control over the order tuples are returned.
  - Fields in ORDER BY clause must also appear in SELECT clause.
- **LIMIT** controls the number of rows returned (good fit w/ORDER BY)
- Can also modify/delete tuple pointed to by a cursor
  - A "non-relational" way to get a handle to a particular tuple

## Database APIs



- A library with database calls (API)
  - special objects/methods
  - passes SQL strings from language, presents **result sets** in a language-friendly way
  - *ODBC* a C/C++ standard started on Windows
  - *JDBC* a Java equivalent
  - Most scripting languages have similar things
    - E.g. in Ruby there's the "pg" gem for Postgres
- ODBC/JDBC try to be DBMS-neutral
  - at least try to hide distinctions across different DBMSs
- Object-Relational Mappings (ORMs)
  - Ruby on Rails, Django, Spring, BackboneORM, etc.
    - Automatically map database rows into PL objects
    - Magic can be great; magic can bite you.
  - This year we won't cover ORMs much – see CS169.

## Summary



- Relational model has **well-defined query semantics**
- SQL provides functionality close to basic relational model  
(some differences in duplicate handling, null values, set operators, ...)
- Typically, many ways to write a query
  - DBMS figures out a fast way to execute a query, regardless of how it is written.



## BACKUP MATERIAL

## Getting Serious



- Two “fancy” queries for different applications
  - Clustering Coefficient for Social Network graphs
  - Medians for “robust” estimates of the central value

## Serious SQL: Social Nets Example



```
-- An undirected friend graph. Store each link once
CREATE TABLE Friends(
  fromID integer,
  toID integer,
  since date,
  PRIMARY KEY (fromID, toID),
  FOREIGN KEY (fromID) REFERENCES Users,
  FOREIGN KEY (toID) REFERENCES Users,
  CHECK (fromID < toID));

-- Return both directions
CREATE VIEW BothFriends AS
SELECT * FROM Friends
UNION ALL
SELECT F.toID AS fromID, F.fromID AS toID, F.since
FROM Friends F;
```

## 6 degrees of friends



```
SELECT F1.fromID, F5.toID
FROM BothFriends F1, BothFriends F2, BothFriends F3,
    BothFriends F4, BothFriends F5
WHERE F1.toID = F2.fromID
AND F2.toID = F3.fromID
AND F3.toID = F4.fromID
AND F4.toID = F5.fromID;
```

## Clustering Coefficient of a Node



$$C_i = 2|\{e_{jk}\}| / k_i(k_i - 1)$$

- where:
  - $k_i$  is the number of neighbors of node  $i$
  - $e_{jk}$  is an edge between nodes  $j$  and  $k$  neighbors of  $i$ , ( $j < k$ ). (A triangle!)
- I.e. Cliquishness: the fraction of your friends that are friends with each other!
- Clustering Coefficient of a graph is the average CC of all nodes.

## In SQL

$$C_i = 2|\{e_{jk}\}| / k_i(k_i - 1)$$




```
CREATE VIEW NEIGHBOR_CNT AS
SELECT
FROM
GROUP
```

```
CREATE VIEW TRIANGLES AS
SELECT
```

```
FROM
WHERE
AND
AND
;
```

## In SQL

$$C_i = 2|\{e_{jk}\}| / k_i(k_i - 1)$$




```
CREATE VIEW NEIGHBOR_CNT AS
SELECT fromID AS nodeID, count(*) AS friend_cnt
FROM BothFriends
GROUP BY nodeID;
```

```
CREATE VIEW TRIANGLES AS
```

```
SELECT
FROM
WHERE
AND
AND
;
```

## In SQL

$$C_i = 2|\{e_{jk}\}| / k_i(k_i - 1)$$



```
CREATE VIEW NEIGHBOR_CNT AS
SELECT fromID AS nodeID, count(*) AS friend_cnt
FROM BothFriends
GROUP BY nodeID;
```

```
CREATE VIEW TRIANGLES AS
```

```
SELECT F1.toID as root, F1.fromID AS friend1,
F2.fromID AS friend2
FROM BothFriends F1, BothFriends F2, Friends F3
WHERE F1.toID = F2.toID /* Both point to root */
AND F1.fromID = F3.fromID /* Same origin as F1 */
AND F3.toID = F2.fromID /* points to origin of F2 */
;
```

## In SQL

$$C_i = 2|\{e_{jk}\}| / k_i(k_i - 1)$$



```
CREATE VIEW NEIGHBOR_EDGE_CNT AS
```

```
SELECT
GROUP
```

```
CREATE VIEW CC_PER_NODE AS
```

```
SELECT
FROM
WHERE
```

```
SELECT AVG(cc) FROM CC_PER_NODE;
```

## In SQL

$$C_i = 2|\{e_{jk}\}| / k_i(k_i - 1)$$



```
CREATE VIEW NEIGHBOR_EDGE_CNT AS
SELECT root, COUNT(*) as cnt FROM TRIANGLES
GROUP BY root;
```

```
CREATE VIEW CC_PER_NODE AS
```

```
SELECT
FROM
WHERE
```

```
SELECT AVG(cc) FROM CC_PER_NODE;
```

## In SQL

$$C_i = 2|\{e_{jk}\}| / k_i(k_i - 1)$$



```
CREATE VIEW NEIGHBOR_EDGE_CNT AS
SELECT root, COUNT(*) as cnt FROM TRIANGLES
GROUP BY root;
```

```
CREATE VIEW CC_PER_NODE AS
```

```
SELECT NE.root, 2.0*NE.cnt /
(N.friend_cnt*(N.friend_cnt-1)) AS CC
FROM NEIGHBOR_EDGE_CNT NE, NEIGHBOR_CNT N
WHERE NE.root = N.nodeID;
```

```
SELECT AVG(cc) FROM CC_PER_NODE;
```

## Median

- Given n values in sorted order, the one at position n/2
  - Assumes an odd # of items
  - For an even #, can take the lower of the middle 2
- A much more "robust" statistic than average
  - Q: Suppose you want the mean to be 1,000,000. What fraction of values do you have to corrupt?
  - Q2: Suppose you want the median to be 1,000,000. Same question.
  - This is called the *breakdown point* of a statistic.
  - Important for dealing with data *outliers*
    - E.g. dirty data
    - Even with real data: "overfitting"



## Median in SQL



```
SELECT c AS median FROM T
WHERE
```



## Median in SQL



```
SELECT c AS median FROM T
WHERE
```



## Median in SQL



```
SELECT c AS median FROM T
WHERE
  (SELECT COUNT(*) FROM T AS T1
   WHERE T1.c < T.c)
```

=



## Median in SQL



```
SELECT c AS median FROM T
WHERE
  (SELECT COUNT(*) FROM T AS T1
   WHERE T1.c < T.c)
=
  (SELECT COUNT(*) FROM T AS T2
   WHERE T2.c > T.c);
```

## Faster Median in SQL



```
SELECT x.c AS median
FROM T x, T y
GROUP BY x.c
HAVING
  SUM(CASE WHEN y.c <= x.c THEN 1 ELSE 0 END)
  >= (COUNT(*)+1)/2
AND
  SUM(CASE WHEN y.c >= x.c THEN 1 ELSE 0 END)
  >= (COUNT(*)/2)+1
```

Why faster?  
Note: handles even # of items!

## Using “Window Functions”



Window functions: an SQL idiom to compute with order.  
<http://www.postgresql.org/docs/9.3/static/tutorial-window.html>

```
CREATE VIEW twocounters AS
(SELECT x,
  ROW_NUMBER() OVER (ORDER BY x ASC) AS RowAsc,
  ROW_NUMBER() OVER (ORDER BY x DESC) AS RowDesc
 FROM numbers
);
```

```
SELECT AVG(x)
FROM twocounters
WHERE RowAsc IN (RowDesc - 1, RowDesc + 1);
```

O(n log n!)  
Note: handles even # of items.

## Notes for Studying



- You'll be responsible for all the constructs we mentioned, except
  - Window functions
  - Programming Language APIs
- In HW3 you may write queries using:
  - Any PostgreSQL features you like
  - Except for callouts to user-defined code (C, Java, Python, R, etc.)